

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
seaborn.set_style('whitegrid')
import warnings
warnings.filterwarnings('ignore')

In [2]: df = pd.read_csv('st.csv')

In [3]: df.head()
```

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob	Fjob	...	internet	romantic	famrel	freetime	goutout	health	absences	G1	G2	G3
0	GP	F	18	U	GT3	A	4	4	at_home	teacher	...	no	no	4	3	4	3	6	5	6	6
1	GP	F	17	U	GT3	T	1	1	at_home	other	...	yes	yes	no	5	3	3	3	4	5	5
2	GP	F	15	U	LE3	T	1	1	at_home	other	...	yes	no	4	3	2	3	10	7	8	10
3	GP	F	15	U	GT3	T	4	2	health	services	...	yes	yes	3	2	2	5	2	15	14	15
4	GP	F	16	U	GT3	T	3	3	other	other	...	no	no	4	3	2	5	4	6	10	10

5 rows × 31 columns

```
In [4]: df.info()


<class 'pandas.core.frame.DataFrame'>
RangeIndex: 395 entries, 0 to 394
Data columns (total 31 columns):
 #   Column      Non-Null Count  Dtype
---  --
 0   school      395 non-null    object
 1   sex         395 non-null    object
 2   age         395 non-null    int64
 3   address     395 non-null    object
 4   famsize     395 non-null    object
 5   Pstatus     395 non-null    object
 6   Medu        395 non-null    int64
 7   Fedu        395 non-null    int64
 8   Mjob        395 non-null    object
 9   Fjob        395 non-null    object
10   reason      395 non-null    object
11   guardian    395 non-null    object
12   traveltime  395 non-null    int64
13   studytime   395 non-null    int64
14   failures    395 non-null    int64
15   schoolsup   395 non-null    object
16   famsup      395 non-null    object
17   paid        395 non-null    object
18   activities  395 non-null    object
19   nursery     395 non-null    object
20   higher      395 non-null    object
21   internet    395 non-null    object
22   romantic    395 non-null    object
23   famrel      395 non-null    int64
24   freetime    395 non-null    int64
25   goutout     395 non-null    int64
26   health      395 non-null    int64
27   absences    395 non-null    int64
28   G1          395 non-null    int64
29   G2          395 non-null    int64
30   G3          395 non-null    int64
dtypes: int64(14), object(17)
memory usage: 95.8+ KB

In [5]: df.describe()
```

	age	Medu	Fedu	traveltime	studytime	failures	famrel	freetime	goutout	health	absences	G1	G2	G3
count	395.000000	395.000000	395.000000	395.000000	395.000000	395.000000	395.000000	395.000000	395.000000	395.000000	395.000000	395.000000	395.000000	395.000000
mean	16.696203	2.749367	2.521519	1.448101	2.035443	0.334177	3.944304	3.235443	3.108861	3.554430	5.709861	10.908961	10.713924	10.415190
std	1.276043	1.004715	1.086201	0.697505	0.839240	0.743651	0.896659	0.998962	1.113278	1.390303	8.003090	3.319195	3.761505	4.581443
min	15.000000	0.000000	0.000000	1.000000	1.000000	0.000000	1.000000	1.000000	1.000000	1.000000	0.000000	3.000000	0.000000	0.000000
25%	16.000000	2.000000	2.000000	1.000000	1.000000	0.000000	4.000000	3.000000	2.000000	3.000000	0.000000	8.000000	9.000000	8.000000
50%	17.000000	3.000000	2.000000	1.000000	2.000000	0.000000	4.000000	3.000000	3.000000	4.000000	4.000000	11.000000	11.000000	11.000000
75%	18.000000	4.000000	3.000000	2.000000	2.000000	0.000000	5.000000	4.000000	4.000000	5.000000	8.000000	13.000000	13.000000	14.000000
max	22.000000	4.000000	4.000000	4.000000	4.000000	3.000000	5.000000	5.000000	5.000000	5.000000	75.000000	19.000000	19.000000	20.000000

```
In [6]: plt.figure(figsize=(12,10))
sns.heatmap(df.isnull(),cmap="magma",char=False)

Out[6]: <AxesSubplot>
```

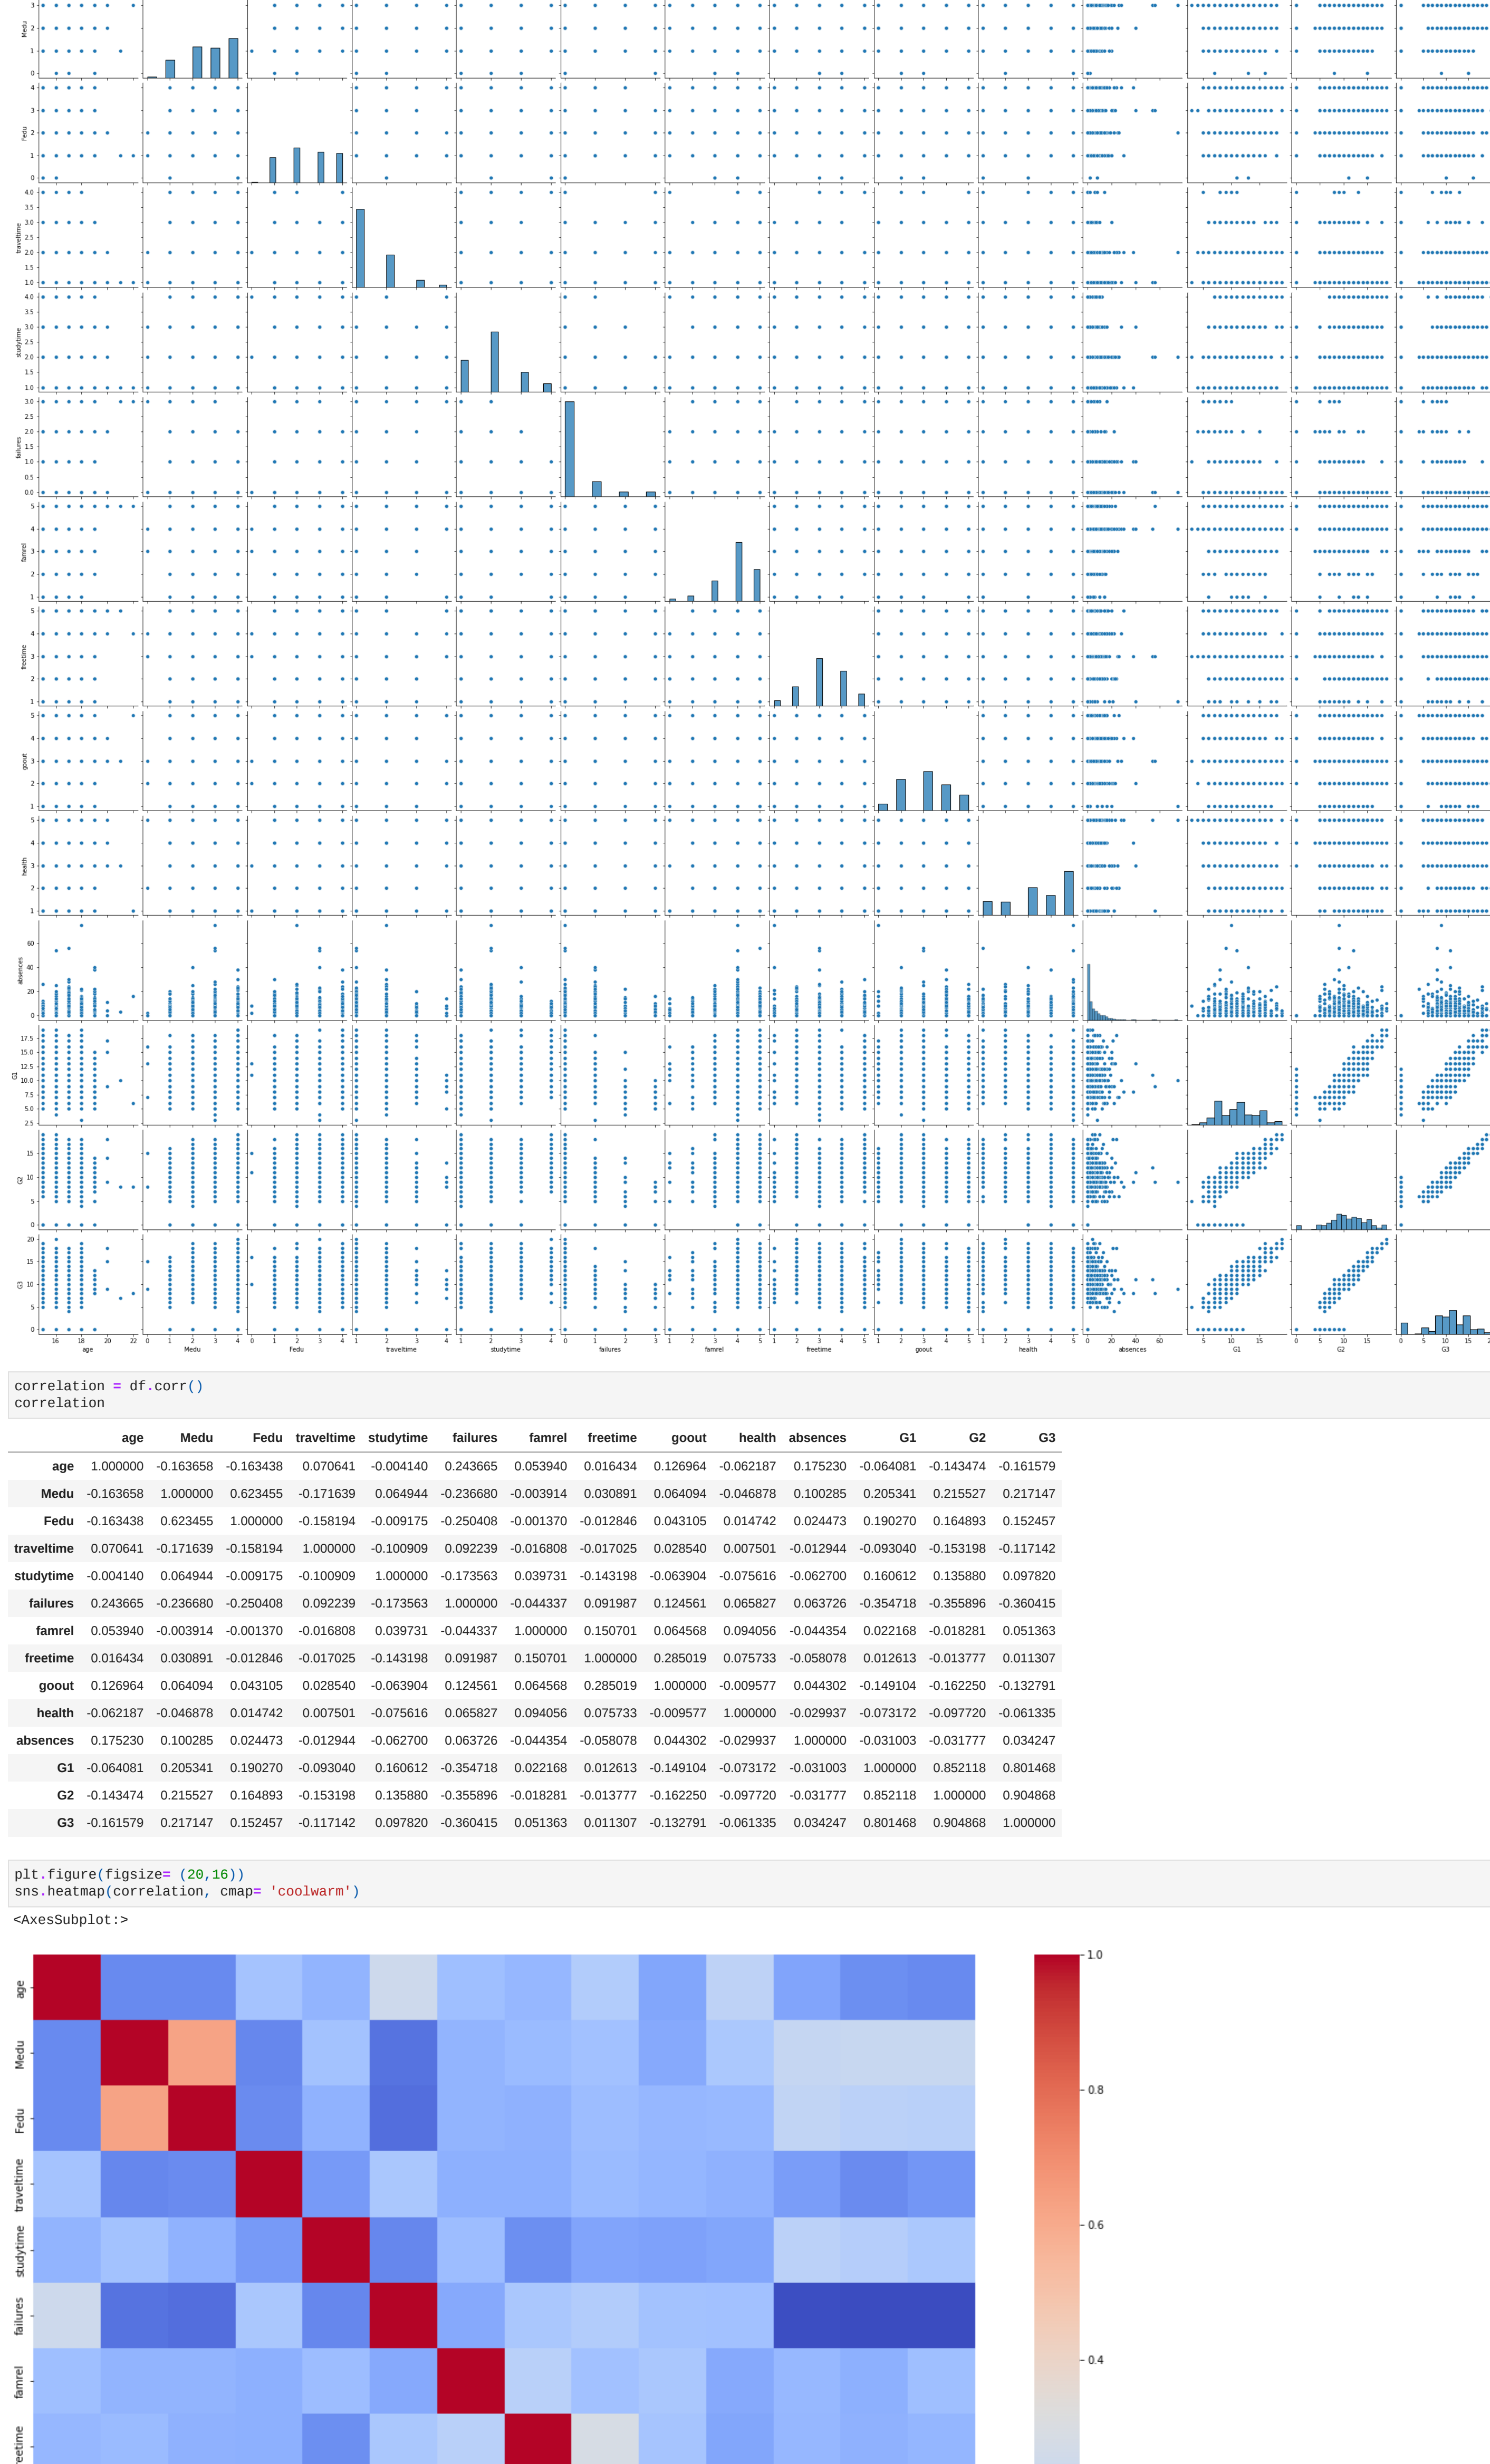


```
In [7]: df.isnull().sum()

Out[7]: school      0
sex            0
age            0
address        0
famsize        0
Pstatus        0
Medu           0
Fedu           0
Mjob           0
Fjob           0
reason         0
guardian       0
traveltime     0
studytime      0
failures       0
schoolsup      0
famsup         0
paid           0
activities     0
nursery        0
higher         0
internet       0
romantic       0
famrel         0
freetime       0
goutout        0
health         0
absences       0
G1             0
G2             0
G3             0
dtype: int64

In [8]: sns.pairplot(df)

Out[8]: <seaborn.axisgrid.PairGrid at 0x291c552ac10>
```



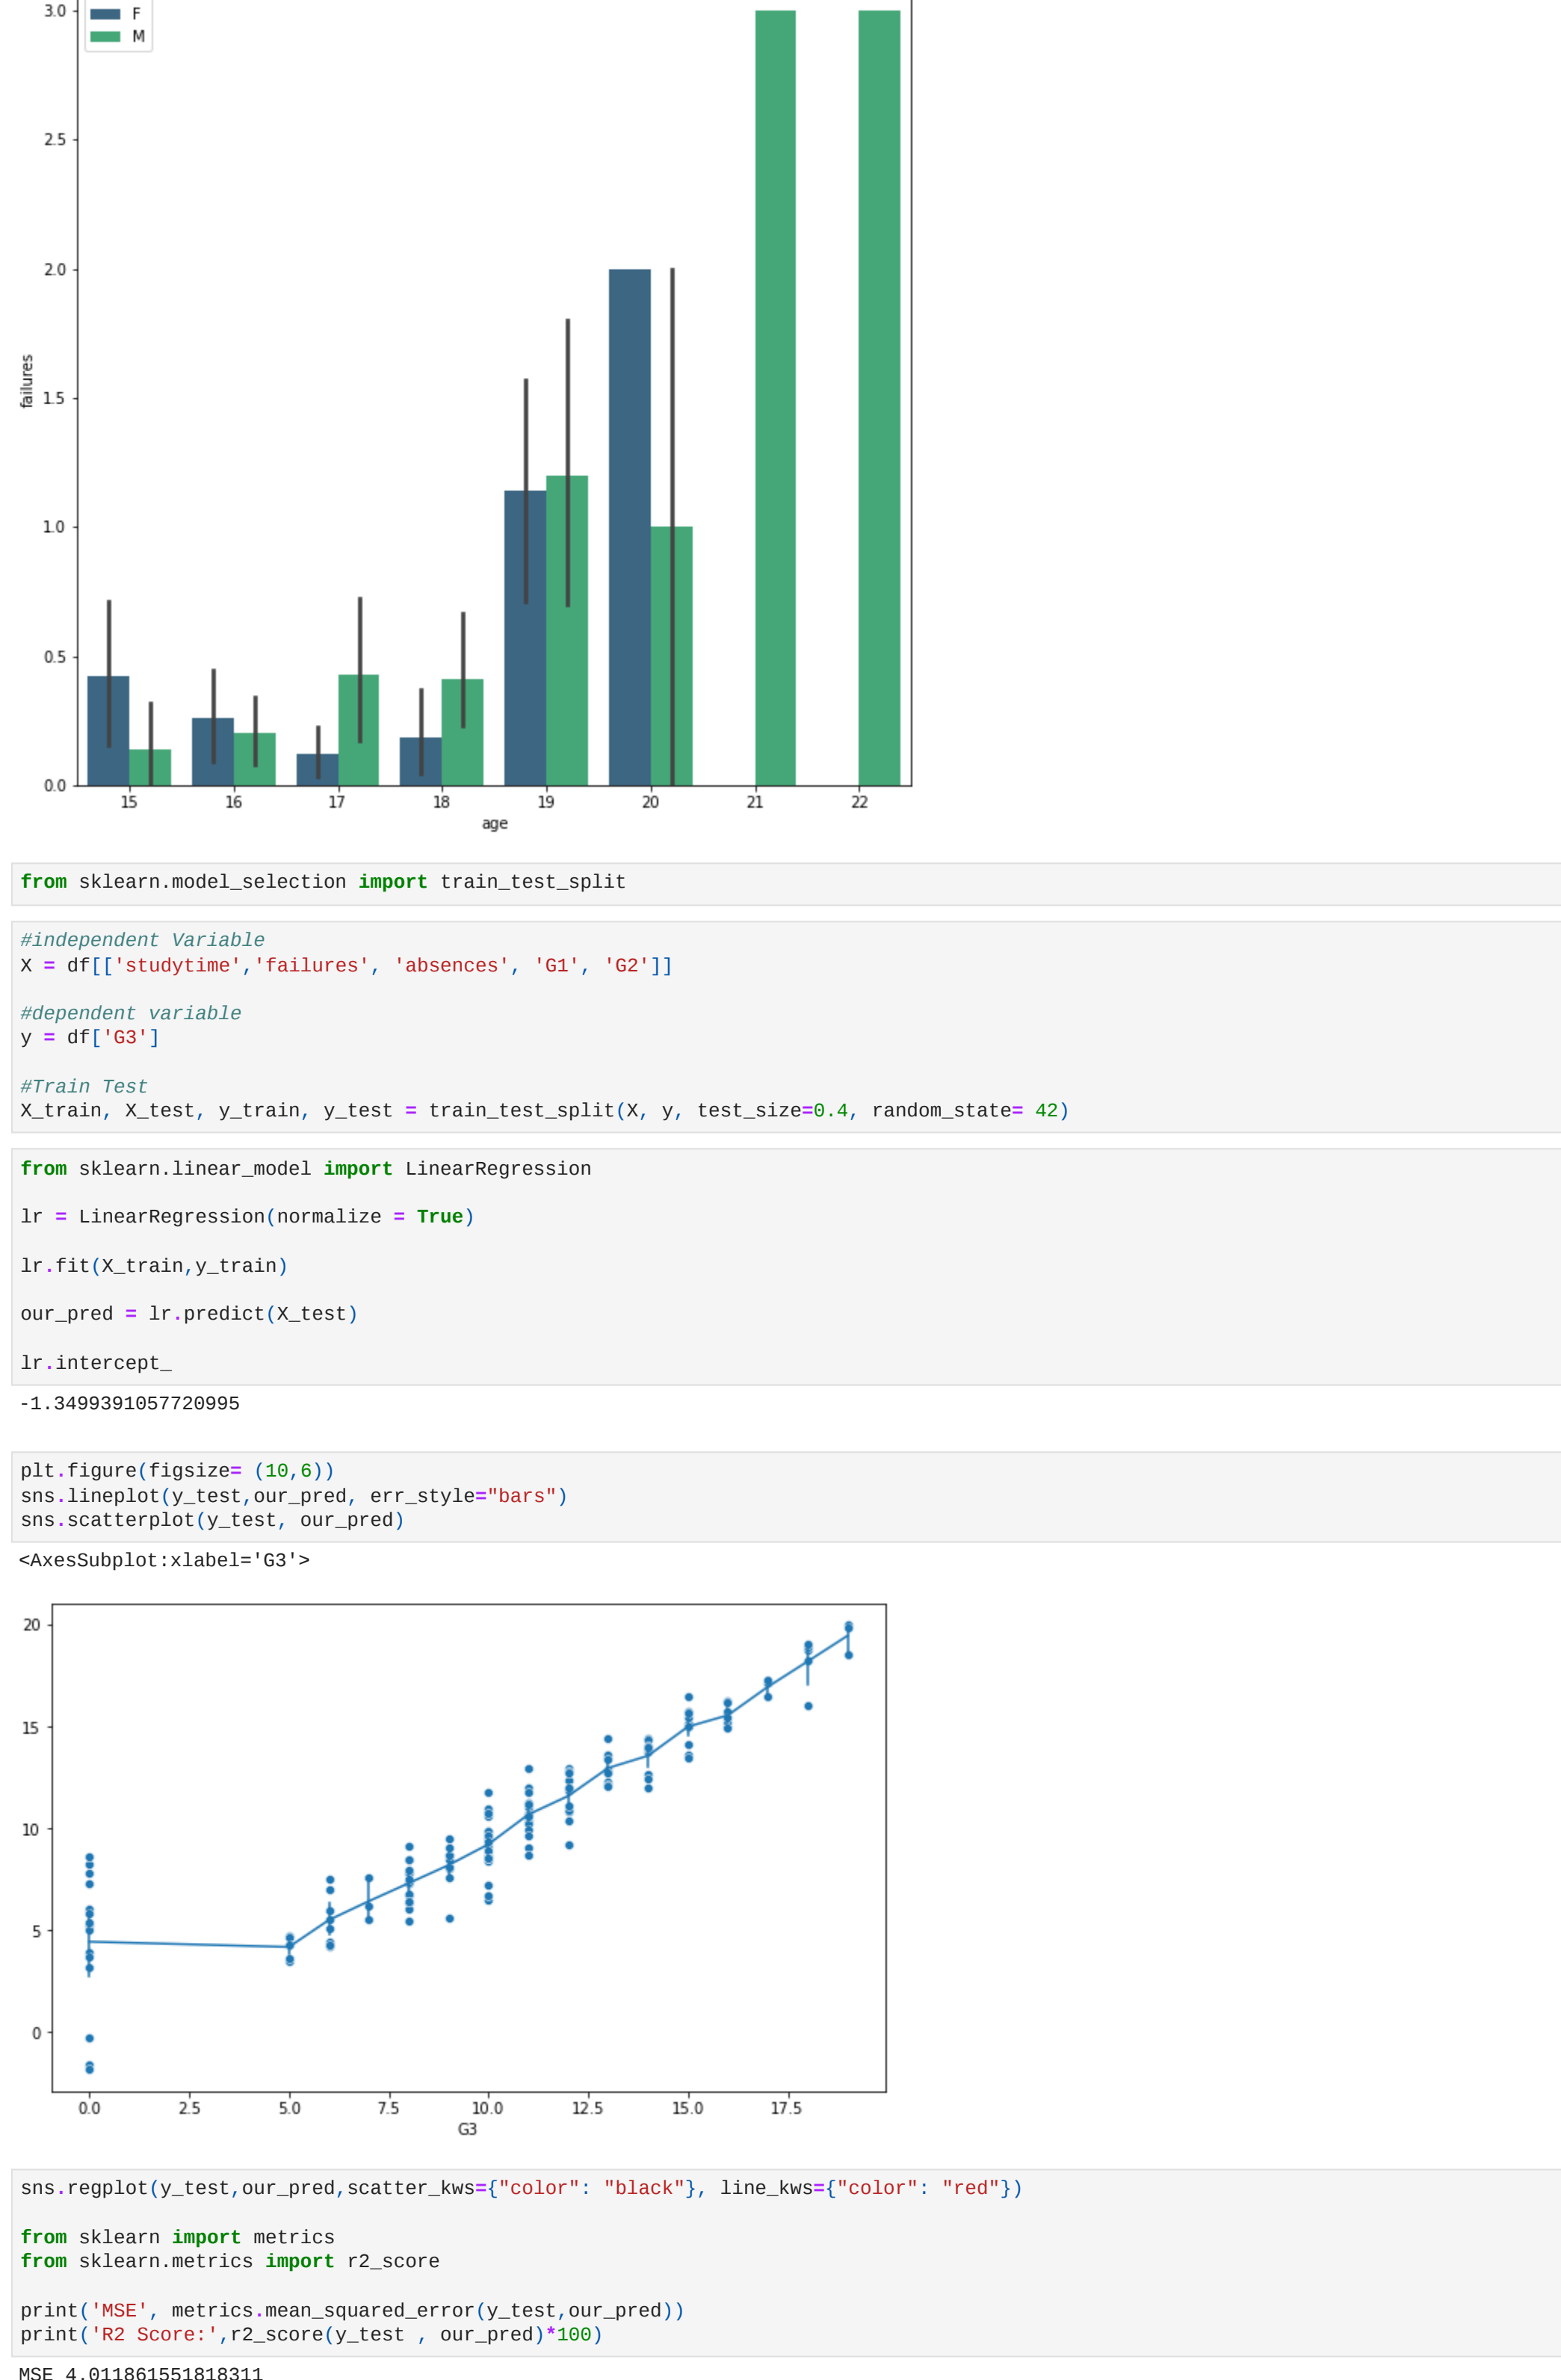
```
In [9]: correlation = df.corr()
correlation

Out[9]:
```

	age	Medu	Fedu	traveltime	studytime	failures	famrel	freetime	goutout	health	absences	G1	G2	G3
age	1.000000	-0.163658	-0.163438	0.070641	-0.004140	0.243665	-0.053940	0.016434	0.126964	-0.062187	0.175230	-0.064081	-0.143474	-0.161579
Medu	-0.163438	1.000000	0.623455	-0.171639	0.064944	-0.236680	-0.003914	0.030891	0.064094	-0.046878	0.100285	0.205341	0.215527	0.217147
Fedu	-0.163438	0.623455	1.000000	-0.158194	-0.009175	-0.250408	-0.001370	-0.012846	0.043105	0.014742	0.024473	0.190270	0.164893	0.152457
traveltime	0.070641	-0.171639	-0.158194	1.000000	-0.100909	0.092239	-0.016808	-0.017025	0.028540	0.007501	-0.012944	-0.083040	-0.153196	-0.117142
studytime	-0.004140	0.064944	-0.009175	-0.100909	1.000000	-0.173563	0.039731	-0.143196	-0.063904	-0.075616	-0.062700	0.106012	0.135880	0.097820
failures	0.243665	-0.236680	-0.250408	0.092239	-0.173563	1.000000	-0.044337	0.091967	0.124561	0.065827	0.063726	-0.354718	-0.358986	-0.360415
famrel	-0.053940	-0.003914	-0.001370	-0.016808	0.039731	-0.044337	1.000000	0.150701	0.064568	0.094056	-0.044354	0.022108	-0.018261	0.051263
freetime	0.016434	0.030891	-0.012846	-0.017025	-0.143196	0.091967	0.150701	1.000000	0.289019	0.075723	0.058078	0.012613	-0.013777	0.011307
goutout	0.126964	0.064094	0.043105	0.028540	-0.063904	0.124561	0.064568	0.289019	1.000000	0.009577	0.044302	-0.014904	-0.162250	-0.132791
health	-0.062187	-0.046878	0.014742	0.007501	-0.075616	0.065827	0.094056	0.075723	0.009577	1.000000	-0.029937	-0.073172	-0.097720	-0.061335
absences	0.175230	0.100285	0.024473	-0.012944	-0.062700	0.063726	-0.044354	-0.058078	0.044302	-0.029937	1.000000	-0.031003	-0.031777	0.034247
G1	-0.064081	0.205341	0.190270	-0.083040	0.106012	-0.354718	0.022108	0.012613	-0.149104	-0.073172	-0.031003	1.000000	0.892118	0.801468
G2	-0.143474	0.215527	0.164893	-0.153196	0.135880	-0.358986	-0.018261	-0.013777	-0.162250	-0.097720	-0.031777	0.892118	1.000000	0.904686
G3	-0.161579	0.217147	0.152457	-0.117142	0.097820	-0.360415	0.051263	0.011307	-0.132791	-0.061335	0.034247	0.801468	0.904686	1.000000

```
In [10]: plt.figure(figsize=(10,10))
sns.heatmap(correlation,cmap= 'coolwarm')

Out[10]: <AxesSubplot>
```



```
In [11]: plt.figure(figsize=(10,10))
sns.barplot(df['age'],df['failures'],hue=df['sex'],palette='viridis')

Out[11]: <AxesSubplot: xlabel='age', ylabel='failures'>
```



```
In [12]: from sklearn.model_selection import train_test_split

In [13]: #Independent Variable
X = df[['studytime','failures', 'absences', 'G1', 'G2']]

#dependent variable
y = df['G3']

#Train Test
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4, random_state= 42)

In [14]: from sklearn.linear_model import LinearRegression

lr = LinearRegression(normalize = True)
lr.fit(X_train,y_train)

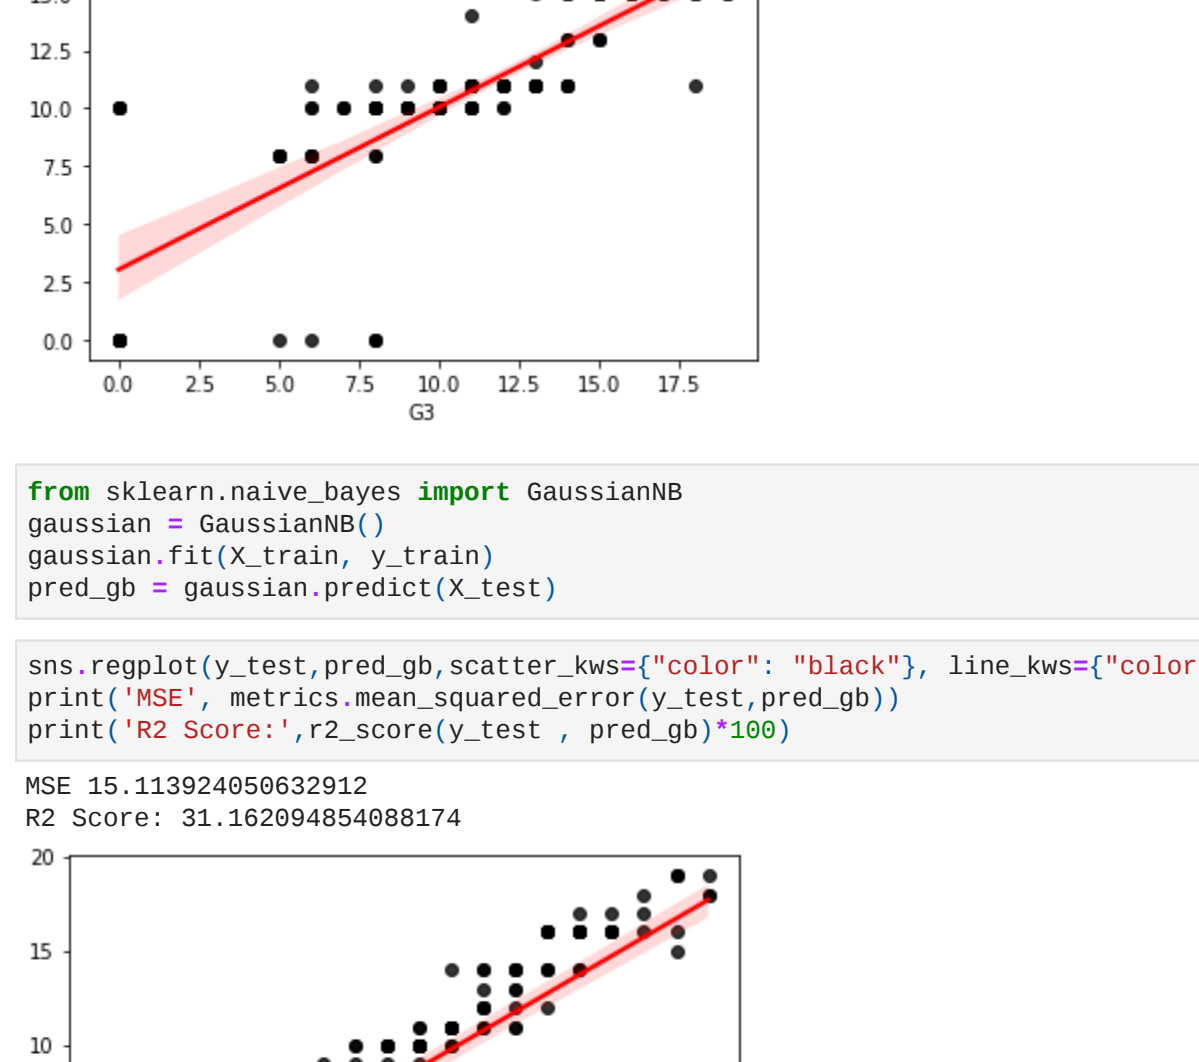
our_pred = lr.predict(X_test)

lr.intercept_

Out[14]: -1.3499391857728995

In [15]: plt.figure(figsize=(10,6))
sns.lineplot(y_test,our_pred, err_style='bars')
sns.scatterplot(y_test, our_pred)

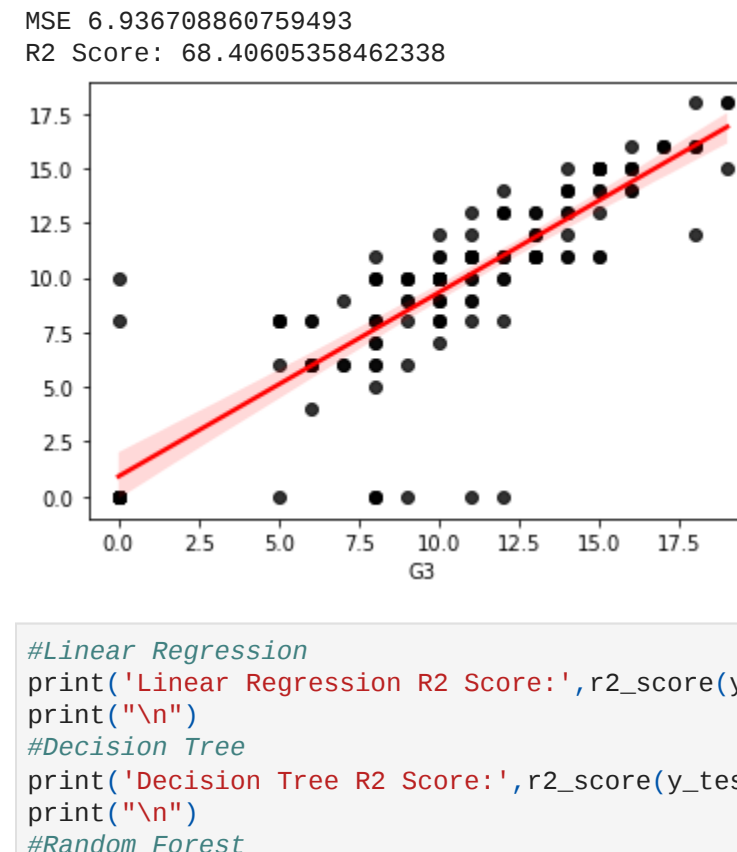
Out[15]: <AxesSubplot: xlabel='G3'>
```



```
In [16]: sns.regplot(y_test,our_pred,scatter_kws={"color": "black"}, line_kws={"color": "red"})

from sklearn.metrics import r2_score
print('MSE', metrics.mean_squared_error(y_test,our_pred))
print('R2 Score:',r2_score(y_test , our_pred)*100)

MSE: 4.611861551818311
R2 Score: 81.72756829811946
```



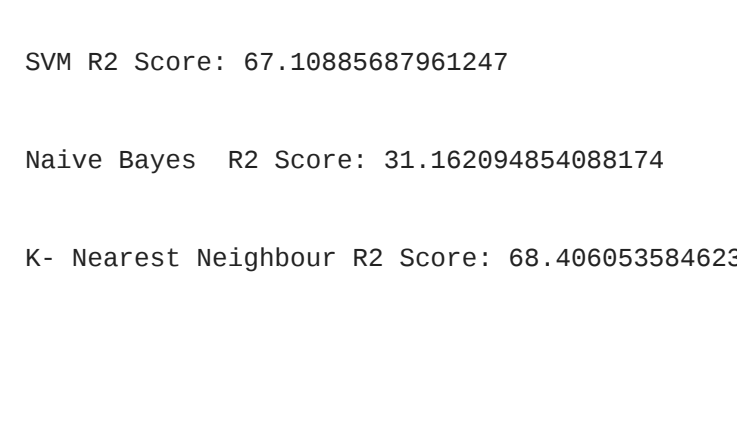
```
In [17]: from sklearn.tree import DecisionTreeRegressor

dt = DecisionTreeRegressor()
dt.fit(X_train,y_train)
pred_t = dt.predict(X_test)

In [18]: sns.regplot(y_test,pred_t,scatter_kws={"color": "black"}, line_kws={"color": "red"})

print('MSE', metrics.mean_squared_error(y_test,pred_t))
print('R2 Score:',r2_score(y_test , pred_t)*100)

MSE: 4.2784810126582276
R2 Score: 80.5132228313918
```



```
In [19]: from sklearn.ensemble import RandomForestRegressor

rf = RandomForestRegressor()
rf.fit(X_train,y_train)
rf_pred = rf.predict(X_test)

In [20]: sns.regplot(y_test,rf_pred,scatter_kws={"color": "black"}, line_kws={"color": "red"})

print('MSE', metrics.mean_squared_error(y_test,rf_pred))
print('R2 Score:',r2_score(y_test , rf_pred)*100)

MSE: 2.564493588196504
R2 Score: 88.36529989551729
```



```
In [21]: from sklearn.svm import SVC

svc = SVC()
svc.fit(X_train, y_train)
pred_svc = svc.predict(X_test)

In [22]: sns.regplot(y_test,pred_svc,scatter_kws={"color": "black"}, line_kws={"color": "red"})

print('MSE', metrics.mean_squared_error(y_test,pred_svc))
print('R2 Score:',r2_score(y_test , pred_svc)*100)

MSE: 7.221519873417724
R2 Score: 67.10885687951247
```



```
In [23]: from sklearn.naive_bayes import GaussianNB

gaussian = GaussianNB()
gaussian.fit(X_train, y_train)
pred_gb = gaussian.predict(X_test)

In [24]: sns.regplot(y_test,pred_gb,scatter_kws={"color": "black"}, line_kws={"color": "red"})

print('MSE', metrics.mean_squared_error(y_test,pred_gb))
print('R2 Score:',r2_score(y_test , pred_gb)*100)

MSE: 15.113924898029212
R2 Score: 31.162094854088174
```



```
In [25]: from sklearn.neighbors import KNeighborsClassifier

knn = KNeighborsClassifier(n_neighbors = 3)
knn.fit(X_train, y_train)
pred_knn = knn.predict(X_test)

In [26]: sns.regplot(y_test,pred_knn,scatter_kws={"color": "black"}, line_kws={"color": "red"})

print('MSE', metrics.mean_squared_error(y_test,pred_knn))
print('R2 Score:',r2_score(y_test , pred_knn)*100)

MSE: 6.236788868799493
R2 Score: 68.4069358462338
```



```
In [27]: #Linear Regression
print('Linear Regression R2 Score:',r2_score(y_test , our_pred)*100)
print("\n")

#Decision Tree
print('Decision Tree R2 Score:',r2_score(y_test , rf_pred)*100)
print("\n")

#Random Forest
print('Random Forest R2 Score:',r2_score(y_test , rf_pred)*100)
print("\n")

#SVM
print('SVM R2 Score:',r2_score(y_test , pred_svc)*100)
print("\n")

#Naive Bayes
print('Naive Bayes R2 Score:',r2_score(y_test , pred_gb)*100)
print("\n")

#K-Nearest Neighbour
print('K- Nearest Neighbour R2 Score:',r2_score(y_test , pred_knn)*100)
print("\n")

Linear Regression R2 Score: 81.72756829811946

Decision Tree R2 Score: 80.5132228313918

Random Forest R2 Score: 88.36529989551729

SVM R2 Score: 67.10885687951247

Naive Bayes R2 Score: 31.162094854088174

K- Nearest Neighbour R2 Score: 68.4069358462338
```